

Programming Assignment 1

Release: February 8

Due: February 21, 23:59

Observing the OS through the /proc file system

The Linux kernel is a collection of data structures. The collective kernel variables define the kernel's perspective of the state of the entire computer system. Each externally invoked function could cause the system state to be changed by having the kernel code change its kernel variables. You can determine the performance and other behaviour of the OS by inspecting those states. In this assignment, we study some aspects of the organization and behaviour of the Linux system by observing values of kernel data structures exposed through the /proc virtual file system.

1. The /proc file system

The /proc filesystem is a virtual filesystem that permits a novel approach for communication between the Linux kernel and user space. In the /proc filesystem, virtual files can be read from or written to as a means of communicating with entities in the kernel. But unlike regular files, the content of these virtual files is dynamically created. You can read a wide range of information on kernel data structures from the /proc filesystem. If you cd into /proc, you will find several files and directories. The directories with numeric names correspond to processes in the system. For example, since the first process created in Linux is the init process with id of 1, the directory named "1" represents the init. It should be noted that files in /proc can be read just like ordinary files, like `fgets()`, `fscanf()`.

The `proc(5)` manual page explains the virtual files and their content available through the /proc filesystem.

```
$ man proc
```

To finish this assignment, you need to check this manual page: figuring out which file can provide the desired information.

2. Problem statement

You are required to write two C programs to report some behaviour of the Linux kernel.

You must read information from the /proc file system. You cannot run any shell command in your program.

Task 1 (40 points)

Write the `info` program, in the source file `info.c`, which displays system information every 3 seconds:

- System uptime, in the format: Uptime: D day(s), H:MM:SS. (10 points)

For example,

```
Uptime: 1 day(s), 7:45:07      /proc/uptime
```

- Memory usage in kB (defined as total memory minus free memory), total memory in kB and the percentage of used memory, in the format: Memory: used kB / total kB (%). The percentage should be displayed with 1 decimal place. (10 points)

For example, `/proc/meminfo`

```
Memory: 902176 kB / 4035964 kB (22.4%)
```

- CPU usage

<https://www.binarytides.com/linux-command-check-memory-usage/>

<https://stackoverflow.com/questions/41224738/how-to-calculate-system-memory-usage-from-proc-meminfo-like-htop>

- Displays the CPU usage of *each* CPU in the 3-second period spent in user mode and kernel mode. (5 points)
- It should not assume the number of CPUs, meaning it should work on machines with different number of CPUs (5 points)
- The usage should be displayed with 1 decimal place, left-padded with space to 5 characters. There should be 2 spaces between each CPU (5 points):
CPU0: 100.0% CPU1: 0.0% CPU2: 0.0% CPU3: 0.0%
- Accepts an optional integer argument to specify the period: `./info N` now shows information every N seconds. You can assume N is an integer. Report error if N is less than or equal to zero. (5 points)

Sample output:

```
$ ./info
(3 seconds later)
Uptime: 1 day(s), 7:59:44
Memory: 902188 kB / 4035964 kB (22.4%)
CPU0: 100.0% CPU1: 0.0% CPU2: 0.0% CPU3: 0.0%

(3 seconds later)
Uptime: 1 day(s), 7:59:47
Memory: 902188 kB / 4035964 kB (22.4%)
CPU0: 100.0% CPU1: 0.0% CPU2: 0.0% CPU3: 0.0%
```

Task 2 (60 points)

Write the `myps` program, in the source file `myps.c`, which displays information for each processes, similar to the utilities `ps` and `top`.

- For each process, display the following information in a row. The rows should be sorted in ascending process ID: (10 points)
 - Process ID (PID)
 - User name of the user who owns this process
 - Total CPU time spent in user and kernel mode
 - Size of virtual memory
 - Command
- The output should be formatted:
 - The PID is left-padded to 7 characters (2 points)
 - The User name is right-padded to 20 characters. (2 points)
 - CPU time is displayed in the format, HH:MM:SS, assumed CPU time never exceeds 100 hours. (2 points)
 - Size of virtual memory is displayed in kB, left-padded to 10 characters. (2 points)
 - Shows the header header in the first row: (2 points)

PID	USER	TIME	VIRT	CMD
-----	------	------	------	-----

- There is a space between each column
- Accepts the `-m` option to sort the processes in descending size of virtual memory. (10 points)
- Accepts the `-p` option to sort the processes in descending CPU time. Assume `-m` and `-p` are not used at the same time. (10 points)
- Accepts the `-u user` option to show processes whose *effective user* is user. (10 points)
 - `-u` should work with `-m` and `-p`. (5 points)
 - Report error if user does not exist. (5 points)

For example,

```
$ ./mys -u no_such_user
invalid user: no_such_user
```

Sample output:

```
$ ./mys -u student -m
  PID USER      TIME    VIRT CMD
16901 student    00:00:00 169840 (sd-pam)
16898 student    00:00:00 18572  systemd
18287 student    00:00:00 13908  sshd
16999 student    00:00:00 13904  sshd
17270 student    00:00:00 13904  sshd
17083 student    00:00:00 13900  sshd
17084 student    00:00:00 10264  bash
17271 student    00:00:00 10240  bash
18288 student    00:00:00 10132  bash
17000 student    00:00:00 5888   sftp-server
18307 student    00:00:00 3100   mys
```

3. Hints

- You can write a simple program containing an infinite loop to introduce CPU load.
- For conversion of user name to user ID, and vice versa, check `<pwd.h>`.
- For directory traversal, check `<dirent.h>`.
- For sorting, check `qsort()` in `<stdlib.h>`.

4. Submission

Submit a zip file containing `info.c`, `mys.c` and a document or Makefile showing how to compile your files.

Your submission will be graded in the machines of CS Lab 2 (`csl2wkXX.cse.ust.hk`, where `XX`=[01..53])

Note: For all of the programming assignments, only the programming language C will be used. Since we are learning the Linux operating system, which is implemented with C, other languages (e.g., C++, Java, Python, and Ruby) simply play no role. Please spend some time refresh the knowledge of C, if you haven't used it for a long time.